

# **MACHINE LEARNING & DATA MINING**



# Recap

Entropy:

$$H(p) = -\mathbb{E}_p \log p$$

Cross entropy:

$$H(p, q) = -\mathbb{E}_p \log q$$

Kullback–Leibler  
divergence:

$$D_{KL}(p \parallel q) = H(p, q) - H(p) = \mathbb{E}_p \left[ \log \frac{p}{q} \right]$$

# Model with hidden parameters

- Model:

$$p(v, h)$$

- Can be defined e.g. in terms on prior on the hidden variables + conditional for the observed ones:

$$p(v | h) \cdot p(h)$$

- I.e. hidden state is ‘the cause’ of the observed state
- Learning with the maximum likelihood method:

$$\log p(v) \rightarrow \max$$

# Learning with max likelihood

- Need to marginalize out the hidden variables:

$$\begin{aligned}\log p(v) &= \log \frac{p(v)p(h | v)}{p(h | v)} = \log \frac{p(h, v)}{p(h | v)} \\ &= \mathbb{E}_{h \sim p(h | v)} \log \frac{p(h, v)}{p(h | v)} \\ &= \mathbb{E}_{h \sim p(h | v)} \log p(h, v) + H(p(h | v))\end{aligned}$$

- This requires the knowledge of the posterior  $p(h | v)$  – typically intractable

# Approximate inference

- We can approximate the posterior  $p(h | v)$  with some  $q(h)$

$$L = \mathbb{E}_{h \sim q(h)} \log p(h, v) + H(q(h))$$

- Let's compare it with the original log-likelihood:

$$\begin{aligned} \log p(v) - L &= \log p(v) - \mathbb{E}_{h \sim q(h)} \log p(h, v) - H(q(h)) \\ &= \mathbb{E}_{h \sim q(h)} [\log p(v) - \log p(h, v) + \log q(h)] \\ &= \mathbb{E}_{h \sim q(h)} [-\log p(h | v) + \log q(h)] = D_{KL}(q(h) || p(h | v)) \end{aligned}$$

# Approximate inference

- We've shown that:

$$L = \mathbb{E}_{h \sim q(h)} \log p(h, v) + H(q(h))$$

$$\log p(v) - L = D_{KL}(q(h) || p(h | v)) \geq 0$$

- This means that  $L$  is the **lower bound** for the true log-likelihood
  - Also called evidence lower bound (**ELBO**) or variational lower bound
- The better  $q$  approximates the posterior — the closer the bound is to the actual log-likelihood
- Instead of maximizing the likelihood we can maximize ELBO!

# ELBO

- Alternative form:

$$\begin{aligned} L &= \mathbb{E}_{h \sim q(h)} \log p(h, v) + H(q(h)) \\ &= \mathbb{E}_{h \sim q(h)} [\log p(v | h) + \log p(h) - \log q(h)] \\ &= \underbrace{\mathbb{E}_{h \sim q(h)} \log p(v | h)}_{\text{Data term}} - \underbrace{D_{KL}(q(h) || p(h))}_{\text{Regularizer}} \end{aligned}$$



# Approximate inference

- There might be different choices for  $q$  depending on the problem
- Some models can give the  $q$  distribution analytically. E.g. expectation maximization (EM) algorithm optimizes ELBO by repeating the following steps:
  - E-step: set  $q$  to equal the posterior precisely (as defined by the current non-optimal model parameters)
  - M-step: completely or partially maximize ELBO with respect to the model parameters (with fixed  $q$ )
- Alternatively we can specify the form of  $q$  such that it is easy to sample from and calculate KL divergence with the prior



# Bayesian NN

- Key idea: treat the weights as hidden parameters
- Define  $q$  in some simple form (e.g. independent normal distributions for each of the weights)
  - simple to sample from
  - calculate KL-divergence analytically
- Define prior on the weights
  - different priors cause different properties of the network
  - e.g. log-uniform prior favors removing noisy weights [1]
- Optimize the ELBO

[1] D. Molchanov, et. al. Variational Dropout Sparsifies Deep Neural Networks, Published in ICML 2017, <https://arxiv.org/abs/1701.05369>

# Bayesian NN

- Incorporates regularization naturally (using prior)
- Results in an ensemble of networks
  - Robust to errors due to data out of the training set domain
- Can estimate uncertainties
- On-line learning possible (using the learnt weights distribution as the new prior)

# Back-propagating through random operations

$$L = \mathbb{E}_{h \sim q(h)} \log p(v | h) - D_{KL}(q(h) || p(h))$$

Optimizing ELBO requires calculating the gradients  
w.r.t. parameters of  $q$  from which we sample

**Q:** How can we achieve this?

# Back-propagating through random operations

$$L = \mathbb{E}_{h \sim q(h)} \log p(v | h) - D_{KL}(q(h) || p(h))$$

Optimizing ELBO requires calculating the gradients  
w.r.t. parameters of  $q$  from which we sample

**Q:** How can we achieve this?

**A:** Reparametrization trick — define the random sampling as sampling from fixed distribution + differentiable transformation

# Reparametrization trick (example)

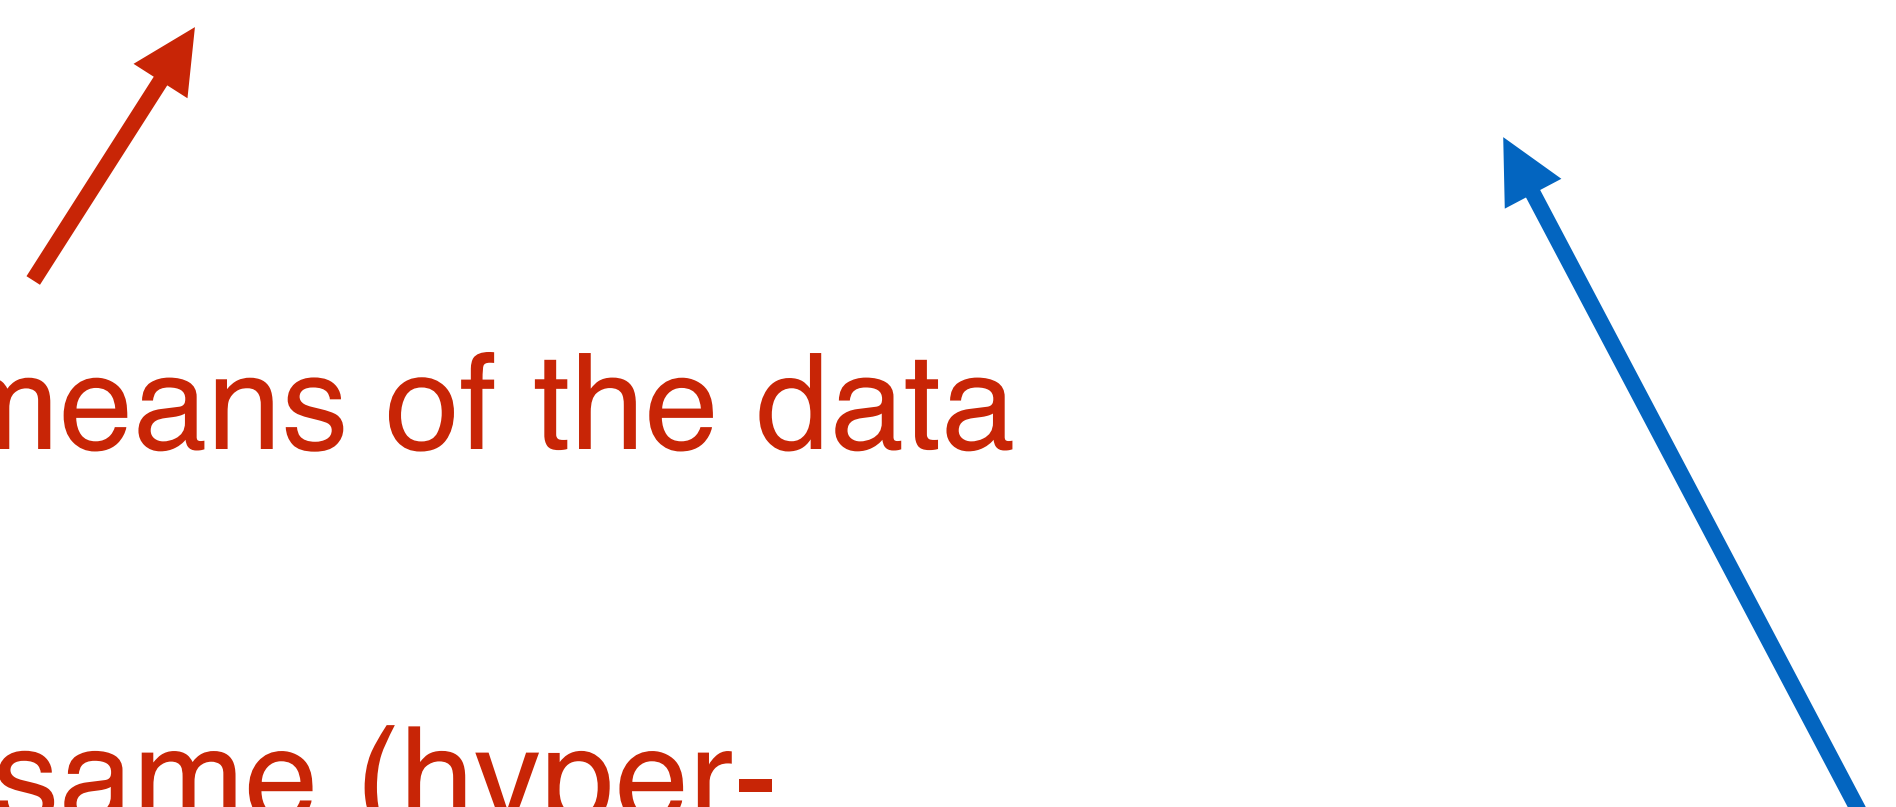
- Say, we want to sample  $y \sim N(\mu, \sigma^2)$  and then optimize some  $f(y)$  wrt to  $\mu$  and  $\sigma$
- We can sample  $z \sim N(0, 1)$
- Then transform it with  $y = \sigma z + \mu$

# Variational Autoencoders (VAE)

- Generative model, learning data distribution  $p(x)$  using a conditional on the hidden variables  $p(x | h)$  and a prior  $p(h)$
- Posterior approximation  $q$  with a neural network – ‘encoder’
  - More precisely:  $q$  is the product of independent normal distributions
  - Encoder neural net inputs a data object
  - and outputs means and variances for these normal distributions
- Conditional  $p(x | h)$  defined as another network – ‘decoder’
- Prior as the product of independent normal distributions with 0 mean and unit variance
- Trained by maximizing the ELBO wrt parameters of encoder and decoder (simultaneously)

# Variational Autoencoders (VAE)

- ELBO:

$$L = \mathbb{E}_{h \sim q(h)} \log p(v | h) - D_{KL}(q(h) || p(h))$$


- Typically decoder predicts means of the data vector
- assuming all variances are same (hyper-parameter)
- This reduces to an MSE loss term
- Variance hyper-parameter controls tradeoff between precision and diversity
- KL-divergence between two normal distributions can be calculated analytically



# Limitations

- When applied to image generation, MSE loss typically results in blurry images

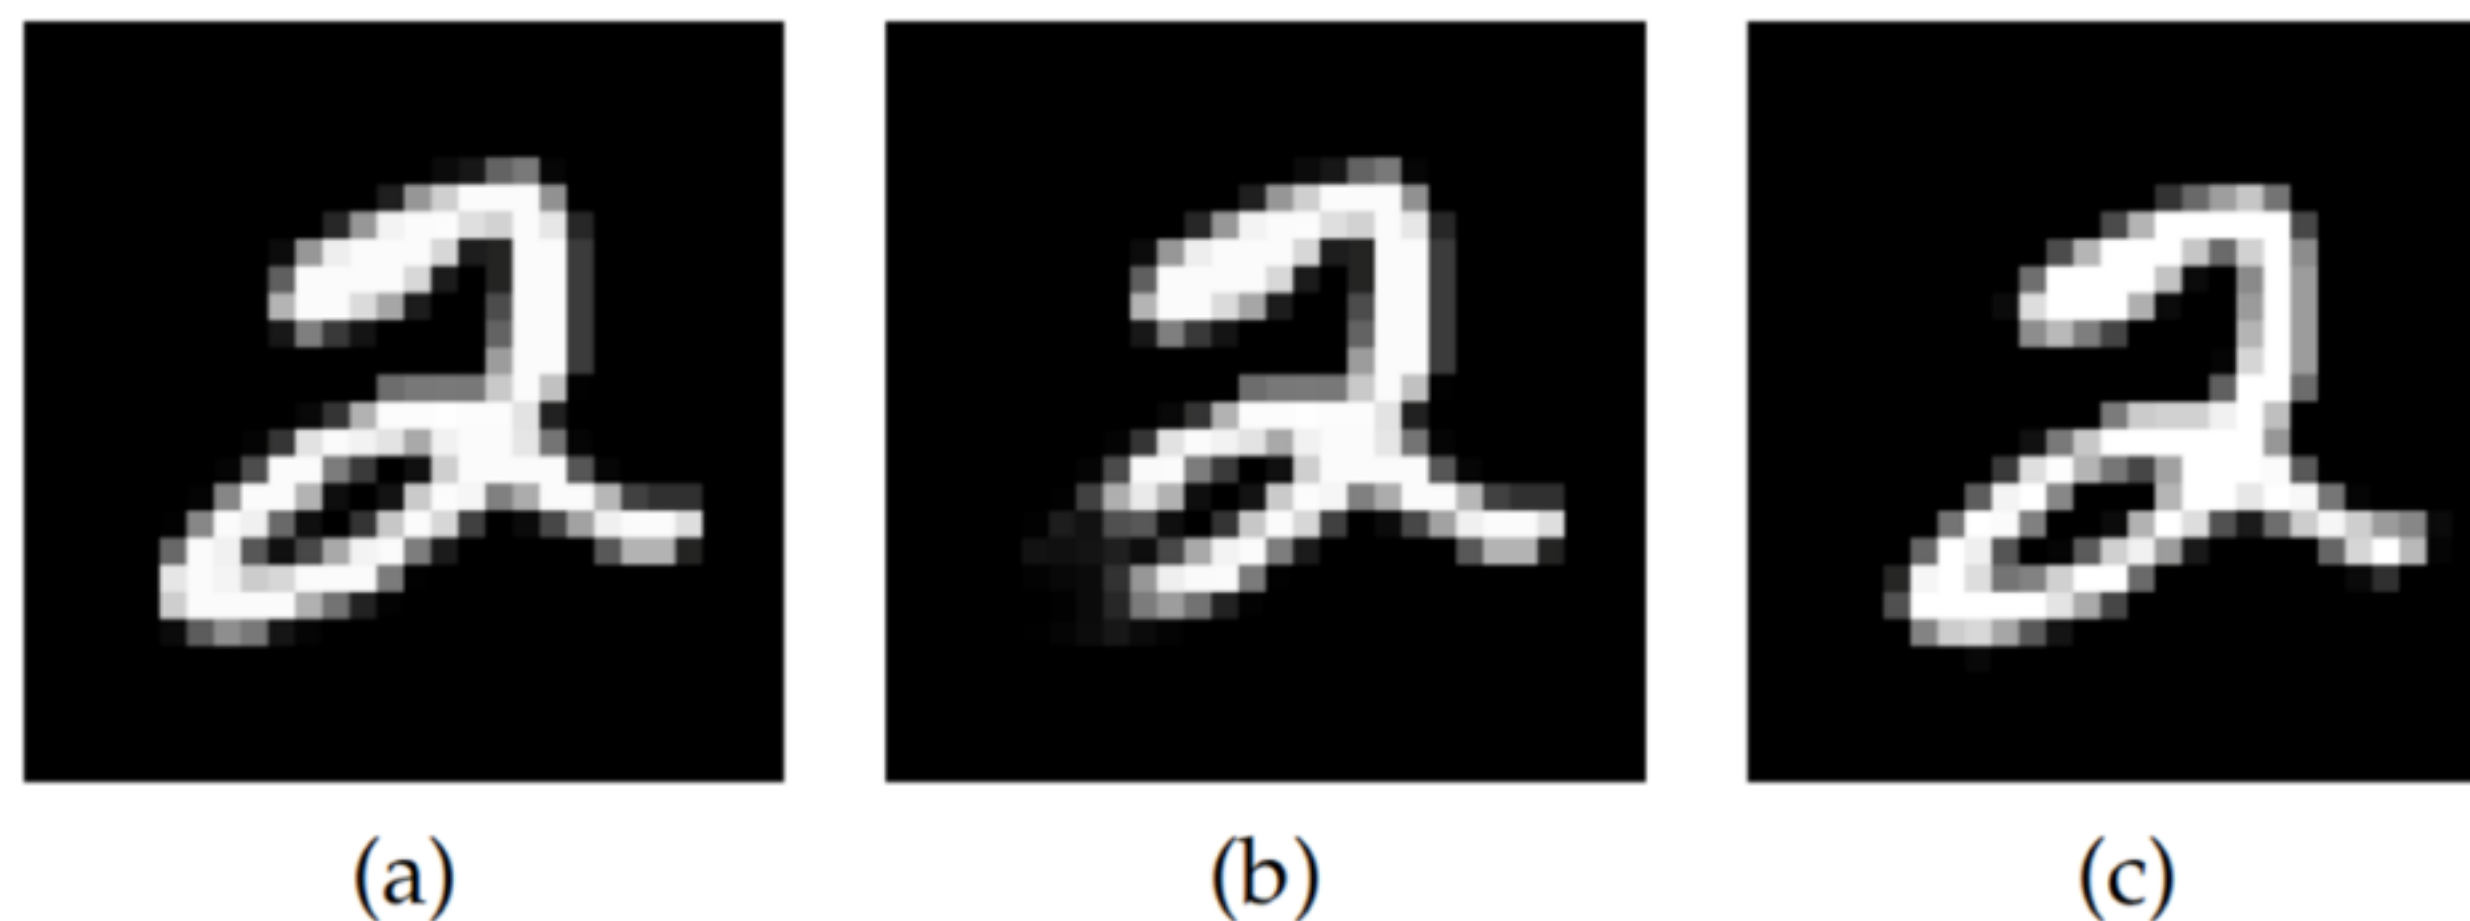


Image (b) — slightly altered image (a), image (c) — image (a) shifted by several pixels.  
Under MSE metric, image (b) is much closer to (a), than (c) to (a).

- MSE loss doesn't reflect our perception of good vs bad image quality